

InfraAgent ReplayOps

Evidence-verified incident triage on AMD

- LangGraph workflow
- Qwen on AMD MI300X via local vLLM
- FastAPI polling API + Hugging Face Space UI

Problem

SRE teams need proof, not another alert stream

- Alert fatigue hides the incident that matters
- Metrics, logs, deploys, traces, and owners are split across tools
- Root-cause notes are hard to audit after the war room

Solution

A bounded agent creates a War Room Packet

- Normalize alert
- Collect deterministic evidence IDs
- Rank root cause with citations
- Generate human-approved runbook
- Verify with an eval scorecard

Architecture

Thin public UI, private AMD inference path

- Hugging Face Space polls Cloudflare Tunnel
- FastAPI runs LangGraph beside vLLM
- vLLM serves Qwen2.5-72B-Instruct on MI300X
- Raw vLLM is never exposed publicly

Judge Proof

The demo shows how every answer was produced

- Runtime truth: live_vllm or honest fallback
- Node trace and tool calls
- Evidence timeline and citation coverage
- War Room Packet with audit seal

Safety

No hidden production claims

- Observability sources are deterministic demo fixtures
- No automatic remediation
- Rollback and scaling require human approval
- Readiness gate blocks submission until live AMD/Qwen proof is visible